

Problem Set 1: Analysis of racial disparities in felony sentencing, Part 1

0. Load packages and imports

```
In [42]: import pandas as pd
import numpy as np
import re

## can add others if you need them

## repeated printouts
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

0.1: Load the data (0 points)

Load the data from `sentencing_asof0405.csv`

- *Notes:* You may receive a warning about mixed data types upon import; feel free to ignore

```
In [43]: df = pd.read_csv("pset1_inputdata/sentencing_asof0405.csv")
```

```
C:\Users\sharp\AppData\Local\Temp\ipykernel_6340\2150987340.py:1: DtypeWarning: Columns (10,11,14,25) have mixed types. Specify dtype option on import or set low_memory=False.
```

```
df = pd.read_csv("pset1_inputdata/sentencing_asof0405.csv")
```

0.2: Print head, dimensions, info (0 points)

```
In [44]: df.head()
df.shape
df.info()
```

Out[44]:

	CASE_ID	CASE_PARTICIPANT_ID	RECEIVED_DATE	OFFENSE_CATEGORY	PRIMARY_C
0	149765331439	175691153649	8/15/1984 12:00:00 AM	PROMIS Conversion	
1	149765331439	175691153649	8/15/1984 12:00:00 AM	PROMIS Conversion	
2	149765331439	175691153649	8/15/1984 12:00:00 AM	PROMIS Conversion	
3	149765331439	175691153649	8/15/1984 12:00:00 AM	PROMIS Conversion	
4	149765331439	175691153649	8/15/1984 12:00:00 AM	PROMIS Conversion	

5 rows × 41 columns

Out[44]: (248146, 41)

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 248146 entries, 0 to 248145
Data columns (total 41 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CASE_ID          248146 non-null   int64  
 1   CASE_PARTICIPANT_ID 248146 non-null   int64  
 2   RECEIVED_DATE     248146 non-null   object  
 3   OFFENSE_CATEGORY  248146 non-null   object  
 4   PRIMARY_CHARGE_FLAG 248146 non-null   bool   
 5   CHARGE_ID         248146 non-null   int64  
 6   CHARGE_VERSION_ID 248146 non-null   int64  
 7   DISPOSITION_CHARGED_OFFENSE_TITLE 248146 non-null   object  
 8   CHARGE_COUNT       248146 non-null   int64  
 9   DISPOSITION_DATE    248146 non-null   object  
 10  DISPOSITION_CHARGED CHAPTER 248146 non-null   object  
 11  DISPOSITION_CHARGED ACT    242771 non-null   object  
 12  DISPOSITION_CHARGED SECTION 242771 non-null   object  
 13  DISPOSITION_CHARGED CLASS  248127 non-null   object  
 14  DISPOSITION_CHARGED AOIC   248122 non-null   object  
 15  CHARGE_DISPOSITION    248146 non-null   object  
 16  CHARGE_DISPOSITION_REASON 904 non-null    object  
 17  SENTENCE_JUDGE        247404 non-null   object  
 18  SENTENCE_COURT_NAME    246761 non-null   object  
 19  SENTENCE_COURT_FACILITY 246216 non-null   object  
 20  SENTENCE_PHASE        248146 non-null   object  
 21  SENTENCE_DATE         248146 non-null   object  
 22  SENTENCE_TYPE         248146 non-null   object  
 23  CURRENT_SENTENCE_FLAG 248146 non-null   bool   
 24  COMMITMENT_TYPE       246464 non-null   object  
 25  COMMITMENT_TERM       246434 non-null   object  
 26  COMMITMENT_UNIT       246434 non-null   object  
 27  LENGTH_OF_CASE_in_Days 229126 non-null   float64 
 28  AGE_AT_INCIDENT      238359 non-null   float64 
 29  RACE                  246879 non-null   object  
 30  GENDER                 247337 non-null   object  
 31  INCIDENT_CITY         228745 non-null   object  
 32  INCIDENT_BEGIN_DATE    239122 non-null   object  
 33  INCIDENT_END_DATE      22008 non-null    object  
 34  LAW_ENFORCEMENT_AGENCY 239405 non-null   object  
 35  LAW_ENFORCEMENT_UNIT    76408 non-null    object  
 36  ARREST_DATE            242981 non-null   object  
 37  FELONY REVIEW_DATE     171907 non-null   object  
 38  FELONY REVIEW RESULT    171907 non-null   object  
 39  ARRAIGNMENT_DATE       229126 non-null   object  
 40  UPDATED_OFFENSE_CATEGORY 248146 non-null   object  

dtypes: bool(2), float64(2), int64(5), object(32)
memory usage: 74.3+ MB

```

Part 1: data cleaning/interpretation

1.1: Understanding the unit of analysis (5 points)

- Print the number of unique values for the following columns. Do so in a way that avoids copying/pasting code for the three:
 - Cases (CASE_ID)
 - People in that case (CASE_PARTICIPANT_ID)
 - Charges (CHARGE_ID)
- Write a couple sentences on the following and show an example of each (e.g., a case involving multiple people):
 - Why there are more unique people than unique cases?
 - Why there are more unique charges than unique people?
- Print the mean and median number of charges per case
- Print the mean and median number of participants per case
- Does the data seem to enable us to follow the same defendant across different cases they're charged in? Write 1 sentence in support of your conclusion.

```
In [45]: # No. of unique values in CASE_ID, CASE_PARTICIPANT_ID, CHARGE_ID, in a way that avoids
unique_counts = df[["CASE_ID", "CASE_PARTICIPANT_ID", "CHARGE_ID"]].nunique()
print(unique_counts)
```

CASE_ID	197519
CASE_PARTICIPANT_ID	211977
CHARGE_ID	229015
dtype:	int64

There are more unique people than unique cases because each unique cases can be committed by one or more unique person. For example, a robbery case involving 2 people.

```
In [46]: # Find a case with multiple participants
participants_counts = df.groupby("CASE_ID")["CASE_PARTICIPANT_ID"].nunique()
multi_participants_ex = participants_counts[participants_counts > 1].index[0]

# Show the participants for that case
participants = df[df["CASE_ID"] == multi_participants_ex]["CASE_PARTICIPANT_ID"].un
print(f"EXAMPLE: CASE_ID {multi_participants_ex} has participants: {participants}")
```

EXAMPLE: CASE_ID 166402790922 has participants: [144234439761 144234534133]

There are more unique charges than unique people because one person can be charged with multiple counts of crime. For example, if someone tried to murder their pregnant wife and robbed them of their possessions then that would be intentional homicide of unborn child, murder, and armed robbery.

```
In [47]: # Find participants with more than one unique charge
participant_charge_counts = df.groupby("CASE_PARTICIPANT_ID")["CHARGE_ID"].nunique()
multi_charge_participant = participant_charge_counts[participant_charge_counts > 1]

# Show the charges for that participant
```

```
charges = df[df["CASE_PARTICIPANT_ID"] == multi_charge_participant]["DISPOSITION_CH"]
print(f"EXAMPLE: CASE_PARTICIPANT_ID {multi_charge_participant} has charges: {charg
```

```
EXAMPLE: CASE_PARTICIPANT_ID 97581722610 has charges: ['INT HOMI OF UNBORN CHILD=38-9-1.2' 'MURDER=720-5\\9-1(A)(1-3)' 'ARMED ROBBERY=720-5\\18-2(A)']
```

```
In [48]: # Mean and median number of charges per case
charges_per_case = df.groupby("CASE_ID")["CHARGE_ID"].nunique()
print("Mean number of charges per case:", charges_per_case.mean())
print("Median number of charges per case:", charges_per_case.median())

# Mean and median number of participants per case
participants_per_case = df.groupby("CASE_ID")["CASE_PARTICIPANT_ID"].nunique()
print("Mean number of participants per case:", participants_per_case.mean())
print("Median number of participants per case:", participants_per_case.median())
```

```
Mean number of charges per case: 1.1594580774507768
```

```
Median number of charges per case: 1.0
```

```
Mean number of participants per case: 1.0731980214561636
```

```
Median number of participants per case: 1.0
```

The data does not seem to enable us to follow the same defendant across different cases because `CASE_PARTICIPANT_ID` is unique to each case-participant combination, not to each defendant. There is no persistent defendant ID across different cases in the visible columns.

1.2.1: Which offense is final? (3 points)

- First, read the data documentation [link](#) and summarize in your own words the differences between `OFFENSE_CATEGORY` and `UPDATED_OFFENSE_CATEGORY`
- Construct an indicator `is_changed_offense` that's True for case-participant-charge observations (rows) where there's a difference between the original charge (offense category) and the most current charge (updated offense category). What are some of the more common changed offenses? (can just print result of `sort_values` based on original offense category)
- Print one example of a changed offense from one of these categories and comment on what the reason may be

`OFFENSE_CATEGORY` is the case's initial broad offense label before charges are finalized, while `UPDATED_OFFENSE_CATEGORY` is the later, revised label recalculated to match the case's primary/most severe charge.

```
In [49]: # Create indicator for changed offense (True if original differs from updated)
is_changed_offense = df["OFFENSE_CATEGORY"].fillna("") != df["UPDATED_OFFENSE_CATEG
# Add to datafram
df["is_changed_offense"] = is_changed_offense
```

```
# Count the charged offenses based on original offense category
charged_offenses_sorted = df[df["is_changed_offense"]].groupby("OFFENSE_CATEGORY").
print(f"Total rows with changed offense: {is_changed_offense.sum()}")
print("\nMost common changed offenses (by original offense category):")
print(charged_offenses_sorted)
```

Total rows with changed offense: 35865

Most common changed offenses (by original offense category):

	OFFENSE_CATEGORY	count
61	PROMIS Conversion	6394
33	DUI	3896
81	UWU - Unlawful Use of Weapon	2155
60	Other Offense	2125
2	Aggravated Battery	1927
..
63	Perjury	4
70	Prostitution	3
22	Benefit Recipient Fraud	2
29	Compelling Gang Membership	2
85	Violate Bail Bond	2

[88 rows x 2 columns]

In [50]:

```
# Print one example of a changed offense from a common category
common_offccat_example = df[df["is_changed_offense"] & (df["OFFENSE_CATEGORY"] != "P
print("Example of a changed offense:")
print(common_offccat_example.to_string(index=False))
```

Example of a changed offense:

OFFENSE_CATEGORY	UPDATED_OFFENSE_CATEGORY	CASE_PARTICIPANT_ID	CHARGE_ID
Attempt Homicide	Domestic Battery	203478864452	89337340966

The change in attempt homicide to domestic battery is based on the intent of the criminal and/or insufficient evidence, where battery is the intent to injury and homicide is intent to kill.

1.2.2: Simplifying the charges (5 points)

Using the field (`UPDATED_OFFENSE_CATEGORY`), create a new field, `simplified_offense_derived`, that simplifies the many offense categories into broader buckets using the following process:

First, combine all offenses beginning with "Aggravated" into a single category without that prefix (e.g., Aggravated Battery and Battery just becomes Battery)

Then:

- Combine all offenses with arson into a single arson category (`Arson`)
- Combine all offenses with homicide into a single homicide category (`Homicide`)

- Combine all offenses with vehicle/vehicular in the name into a single vehicle category (`Vehicle-related`)
- Combine all offenses with battery in the name into a single battery category (`Battery`)

Try to do so efficiently (e.g., write a function and apply to a column, rather than edit the variable repeatedly in separate line for each recoded offense)

Print the difference between the # of unique offenses in the original `UPDATED_OFFENSE_CATEGORY` field and the # of unique offenses in your new `simplified_offense_derived` field

```
In [51]: # Build a function that first identifies aggravated charges, removes the prefix, and applies to a column
def simplify_offense_aggravated(offense):
    if pd.isna(offense):
        return np.nan
    offense_lower = str(offense).lower().strip()
    # Remove Leading "Aggravated" (case-insensitive)
    if offense_lower.startswith("aggravated "):
        offense_lower = offense_lower[len("aggravated "):].strip()
    if "arson" in offense_lower:
        return "Arson"
    elif "homicide" in offense_lower:
        return "Homicide"
    elif "vehic" in offense_lower:
        return "Vehicle-related"
    elif "battery" in offense_lower:
        return "Battery"
    else:
        return offense_lower

# Applying to a column
df["simplified_offense_derived"] = df["UPDATED_OFFENSE_CATEGORY"].apply(simplify_of

# Show the difference between the number of unique offenses before and after simplification
original_unique = df["UPDATED_OFFENSE_CATEGORY"].nunique(dropna=False)
simplified_unique = df["simplified_offense_derived"].nunique(dropna=False)
print(f"Unique offenses before simplification: {original_unique}")
print(f"Unique offenses after simplification: {simplified_unique}")
print(f"Difference: {original_unique - simplified_unique}")
```

Unique offenses before simplification: 79

Unique offenses after simplification: 65

Difference: 14

1.3: Cleaning additional variables (10 points)

Clean the following variables; make sure to retain the original variable in data and use the derived suffix so it's easier to pull these cleaned out variables later (e.g., `age_derived`) to indicate this was a transformation

- Race: create True/false indicators for `is_black_derived` (Black only or mixed race with hispanic), Non-Black Hispanic, so either hispanic alone or white hispanic (`is_hisp_derived`), White non-hispanic (`is_white_derived`), or none of the above (`is_othereth_derived`)
- Gender: create a boolean true/false indicator for `is_male_derived` (false is female, unknown, or other)
- Age at incident: you notice outliers like 130-year olds. Winsorsize the top 0.01% of values to be equal to the 99.99th percentile value pre-winsorization. Call this `age_derived`
- Create `sentenceymd_derived` that's a version of `SENTENCE_DATE` converted to datetime format. Also create a rounded version, `sentenceym_derived`, that's rounded down to the first day of the month (e.g., `1/5/2016` would become `1/1/2016` and `3/27/2018` would become `3/1/2018`)
 - Hint: all timestamps are midnight so u can strip in conversion. For full credit, before converting, you notice that some of the years have been mistranscribed (e.g., 291X or 221X instead of 201X). Programatically fix those (eg 2914 -> 2014). Even after cleaning, there will still be some that are after the year 2021 that we'll filter out later. For partial credit, you can ignore the timestamps that cause errors and set `errors = "coerce"` within `pd.to_datetime()` to allow the conversion to proceed.
- Sentencing judge: create an identifier (`judgeid_derived`) for each unique judge (`SENTENCE_JUDGE`) structured as `judge_1, judge_2...,` with the order determined by sorting the judges (will sort on fname then last). When finding unique judges, there are various duplicates we could weed out --- for now, just focus on (1) the different iterations of Doug/Douglas Simpson, (2) the different iterations of Shelley Sutker (who appears both with her maiden name and her hyphenated married name).
 - Hint: due to mixed types, you may need to cast the `SENTENCE_JUDGE` var to a diff type to sort

After finishing, print a random sample of 10 rows (data.sample(n = 10)) with the original and cleaned columns for the relevant variables to validate your work

```
In [52]: # Print all unique values of RACE column for inspection
print("All unique values in RACE column:")
print(sorted(df["RACE"].dropna().unique()))
```

```
All unique values in RACE column:
['ASIAN', 'American Indian', 'Asian', 'Biracial', 'Black', 'HISPANIC', 'Unknown', 'White', 'White [Hispanic or Latino]', 'White/Black [Hispanic or Latino]']
```

```
In [53]: # Race Indicators
race_str = df["RACE"].str.upper().fillna("")
df["is_black_derived"] = (
```

```

        race_str.isin(["BLACK"]) |
        (race_str.str.contains("BLACK") & race_str.str.contains("HISPANIC"))
    )

df["is_hisp_derived"] = (
    race_str.str.contains("HISPANIC") | race_str.str.contains("LATINO")
) & ~df["is_black_derived"]

df["is_white_derived"] = (
    race_str.str.contains("WHITE") &
    ~race_str.str.contains("HISPANIC") &
    ~race_str.str.contains("LATINO") &
    ~df["is_black_derived"]
)

df["is_othereth_derived"] = ~(
    df["is_black_derived"] | df["is_hisp_derived"] | df["is_white_derived"]
)

```

In [54]: # Output code for race indicators

```

race_cols = ["RACE", "is_black_derived", "is_hisp_derived", "is_white_derived", "is_othereth_derived"]

def show_race_examples(df):
    # Find one example for each derived group
    examples = []
    ex_black = df[df["is_black_derived"]].head(1)
    if not ex_black.empty:
        examples.append(ex_black)
    ex_hisp = df[df["is_hisp_derived"]].head(1)
    if not ex_hisp.empty:
        examples.append(ex_hisp)
    ex_white = df[df["is_white_derived"]].head(1)
    if not ex_white.empty:
        examples.append(ex_white)
    ex_other = df[df["is_othereth_derived"]].head(1)
    if not ex_other.empty:
        examples.append(ex_other)
    if examples:
        display(pd.concat(examples)[race_cols])
    else:
        print("No examples found for the derived race indicators.")

print("Targeted examples for each race indicator:")
show_race_examples(df)

print("\nRace indicator value counts:")
for col in race_cols[1:]:
    print(df[col].value_counts())
    print()

```

Targeted examples for each race indicator:

	RACE	is_black_derived	is_hisp_derived	is_white_derived	is_othereth_derived
0	Black	True	False	False	False
18	White [Hispanic or Latino]	False	True	False	False
26	White	False	False	True	False
41	NaN	False	False	False	True

Race indicator value counts:

```
is_black_derived
True      165661
False     82485
Name: count, dtype: int64
```

```
is_hisp_derived
False     204325
True      43821
Name: count, dtype: int64
```

```
is_white_derived
False     212785
True      35361
Name: count, dtype: int64
```

```
is_othereth_derived
False    244843
True      3303
Name: count, dtype: int64
```

```
In [ ]: print(df["GENDER"].str.upper().unique())
```

```
# Gender Indicator: only true if MALE
gender_upper = df["GENDER"].str.upper().fillna("")
df["is_male_derived"] = gender_upper.isin(["MALE"])
print(df["is_male_derived"].value_counts())
```

```
['MALE' 'FEMALE' nan 'MALE NAME, NO GENDER GIVEN' 'UNKNOWN GENDER'
 'UNKNOWN']
is_male_derived
True      217610
False     30536
Name: count, dtype: int64
```

```
In [56]: # Age at Incident Winsorization (Basically capping the 99.99%)
age_col = "AGE_AT INCIDENT"
percentile_9999 = df[age_col].quantile(0.9999)
df["age_derived"] = df[age_col].clip(upper=percentile_9999)
print(df["age_derived"].describe())
```

```
count    238359.000000
mean      32.302611
std       11.779161
min      17.000000
25%     23.000000
50%     29.000000
75%     40.000000
max      81.000000
Name: age_derived, dtype: float64
```

```
In [57]: # Sentence Date Reformating to datetime
```

```
# Function to fix mistranscribed years
def fix_year(date_str):
    if pd.isna(date_str):
        return date_str

    s = str(date_str).strip()
    if not s:
        return s

    parts = s.split()
    date_part = parts[0]
    rest = " ".join(parts[1:])

    # Only handle slash dates like MM/DD/YYYY
    if "/" not in date_part:
        return s

    d = date_part.split("/")
    if len(d) != 3:
        return s

    mm, dd, yyyy = d[0], d[1], d[2]

    # Fix 4-digit years where digits are mistranscribed
    if yyyy.isdigit() and len(yyyy) == 4:
        year_int = int(yyyy)
        if year_int > 2025 and year_int < 9999:
            yyyy = "20" + yyyy[2:]

    fixed_date = f"{mm}/{dd}/{yyyy}"
    return f"{fixed_date} {rest}".strip()

test_dates = [
    "05/01/2914 12:00:00 AM",
    "12/31/2218 12:00:00 AM",
    "01/01/2119 12:00:00 AM",
    "06/15/2020 12:00:00 AM",
    "07/20/2015 12:00:00 AM",
    "03/10/2512 12:00:00 AM",
    "03/10/3212 12:00:00 AM",
    None,
    "random text"
]
```

```

# Print before and after for each test case
for d in test_dates:
    print(f"Original: {d} --> Fixed: {fix_year(d)}")

# Apply the fix function and convert to datetime
df["SENTENCE_DATE_cleaned"] = df["SENTENCE_DATE"].apply(fix_year)
df["sentenceymd_derived"] = pd.to_datetime(df["SENTENCE_DATE_cleaned"], format="%m/%d/%Y")

# Round down to first day of month
df["sentenceym_derived"] = df["sentenceymd_derived"].dt.to_period("M").dt.start_time

# Drop column no longer needed since transferred to sentenceymd_derived
df.drop(columns=["SENTENCE_DATE_cleaned"], inplace=True)

```

```

Original: 05/01/2914 12:00:00 AM --> Fixed: 05/01/2014 12:00:00 AM
Original: 12/31/2218 12:00:00 AM --> Fixed: 12/31/2018 12:00:00 AM
Original: 01/01/2119 12:00:00 AM --> Fixed: 01/01/2019 12:00:00 AM
Original: 06/15/2020 12:00:00 AM --> Fixed: 06/15/2020 12:00:00 AM
Original: 07/20/2015 12:00:00 AM --> Fixed: 07/20/2015 12:00:00 AM
Original: 03/10/2512 12:00:00 AM --> Fixed: 03/10/2012 12:00:00 AM
Original: 03/10/3212 12:00:00 AM --> Fixed: 03/10/2012 12:00:00 AM
Original: None --> Fixed: None
Original: random text --> Fixed: random text

```

In [58]: # Judge ID Creation

```

# Function to standardize judge names before creating mapping
def standardize_judge(j):
    if pd.isna(j):
        return j

    j = str(j).strip().lower()

    if j == "":
        return pd.NA

    if ("simpson" in j) and ("doug" in j or "douglas" in j):
        return "Douglas Simpson"
    if ("sutker" in j) and ("shelley" in j):
        return "Shelley Sutker"

    return j.title() # Capitalize names for consistency

# Apply standardization
df["judge_cleaned"] = df["SENTENCE_JUDGE"].astype(str).apply(standardize_judge)

# Get unique judges, remove missing values and null values, and sort them
unique_judges = df["judge_cleaned"].dropna().unique()
unique_judges_sorted = sorted(unique_judges)

# Create a mapping from judge name to ID
judge_mapping = {judge: f"judge_{i+1}" for i, judge in enumerate(unique_judges_sorted)}
df["judgeid_derived"] = df["judge_cleaned"].map(judge_mapping)

# Dropped column due to mapping on to judgeid_derived

```

```

df.drop(columns=["judge_cleaned"], inplace=True)

# Show mapping for Doug/Douglas Simpson and Shelley Sutker variations
simpson_mask = df["SENTENCE_JUDGE"].astype(str).str.contains("Simpson")
sutker_mask = df["SENTENCE_JUDGE"].astype(str).str.contains("Sutker")

print("Doug/Douglas Simpson mapping:")
print(df.loc[simpson_mask, ["SENTENCE_JUDGE", "judgeid_derived"]].drop_duplicates())

print("\nShelley Sutker variations mapping:")
print(df.loc[sutker_mask, ["SENTENCE_JUDGE", "judgeid_derived"]].drop_duplicates())

print("\nGeneral sample of judge mappings:")
print(df[["SENTENCE_JUDGE", "judgeid_derived"]].drop_duplicates().head(10))

```

Doug/Douglas Simpson mapping:

	SENTENCE_JUDGE	judgeid_derived
1115	Doug Simpson	judge_71
1621	Douglas J Simpson	judge_71

Shelley Sutker variations mapping:

	SENTENCE_JUDGE	judgeid_derived
131	Shelley Sutker	judge_281
142	Shelley Sutker-Dermer	judge_281

General sample of judge mappings:

	SENTENCE_JUDGE	judgeid_derived
0	John Mannion	judge_140
4	Clayton Jay Crane	judge_41
11	James L Rhodes	judge_114
12	Thomas V Gainer	judge_312
18	Kay M Hanlon	judge_163
19	William J Kunkle	judge_333
20	Evelyn B Clay	judge_84
26	Timothy Joseph Joyce	judge_316
27	Steven J Goebel	judge_289
29	Carol M Howard	judge_32

In [61]: # Group the columns that are needed for validation

```

validation_col = ["RACE", "is_black_derived", "is_hisp_derived", "is_white_derived",
                  "GENDER", "is_male_derived",
                  "AGE_AT INCIDENT", "age_derived",
                  "SENTENCE_DATE", "sentenceymd_derived", "sentenceym_derived",
                  "SENTENCE_JUDGE", "judgeid_derived"]

print("Random sample of 10 rows with original vs cleaned columns")
df.sample(n=10)[validation_col]

```

Random sample of 10 rows with original vs cleaned columns

Out[61]:

	RACE	is_black_derived	is_hisp_derived	is_white_derived	is_othereth_derived	G
26970	White	False	False	True	False	
212818	Black	True	False	False	False	
228985	Black	True	False	False	False	
59326	White [Hispanic or Latino]	False	True	False	False	
86994	Black	True	False	False	False	
1831	Black	True	False	False	False	
197013	Black	True	False	False	False	
75650	White	False	False	True	False	
79633	Black	True	False	False	False	
184935	White [Hispanic or Latino]	False	True	False	False	

1.4: Subsetting rows to analytic dataset (5 points)

You decide based on the above to simplify things in the following ways:

- Subset to cases where only one participant is charged, since cases with >1 participant might have complications like plea bargains/informing from other participants affecting the sentencing of the focal participant
- To go from a participant-case level dataset, where each participant is repeated across charges tied to the case, to a participant-level dataset, where each participant has one charge, subset to a participant's primary charge and their current sentence (`PRIMARY_CHARGE_FLAG` is True and `CURRENT_SENTENCE_FLAG` is True). Double check that this worked by confirming there are no longer multiple charges for the same case-participant

- Filter out observations where judge is nan or nonsensical (indicated by is.null or equal to FLOOD)
- Subset to sentencing date between 01-01-2012 and 04-05-2021 (inclusive)

After completing these steps, print the number of rows in the data

```
In [60]: print("Original DataFrame shape:", df.shape, "\n")
# Subset cases where only 1 participant is charged
single_participant_cases = participants_counts[participants_counts == 1].index
df_analytic = df[df["CASE_ID"].isin(single_participant_cases)].copy()
print("After subsetting to single-participant cases.", df_analytic.shape, "\n")

# Reduce to a participant-level dataset with only primary charges and current sentences
focal = df_analytic["PRIMARY_CHARGE_FLAG"] & df_analytic["CURRENT_SENTENCE_FLAG"]
df_analytic = df_analytic.loc[focal].copy()
print("After subsetting to primary charges with current sentences:", df_analytic.shape)

# Confirm no multiple charges for the same case-participant
charges_per_case_participant = df_analytic.groupby(["CASE_ID", "CASE_PARTICIPANT_ID"])
print("Max unique charges per case-participant after subsetting:", charges_per_case_participant.max())
print("Min unique charges per case-participant after subsetting:", charges_per_case_participant.min())

# Filter out observations where judge is nan or nonsensical
j = df_analytic["SENTENCE_JUDGE"]
judge_ok = j.notna() & j.astype(str).str.strip().ne("") & j.astype(str).str.strip().ne("FLOOD")
df_analytic = df_analytic.loc[judge_ok].copy()

print("Drop n/a judges:", df_analytic.shape, "\n")

# Subset to sentencing date between 01-01-2012 and 04-05-2021 inclusive
start_date = pd.Timestamp("2012-01-01")
end_date = pd.Timestamp("2021-04-05")

date_ok = df_analytic["sentenceymd_derived"].between(start_date, end_date, inclusive="both")
df_analytic = df_analytic.loc[date_ok].copy()

# Print number of rows after subsetting
print("Number of rows in analytic dataset after subsetting:", len(df_analytic))

# Sanity Check
print("\nSanity Check of Analytic Dataset:")

print("Unique CASE_PARTICIPANT_ID in analytic dataset:", df_analytic["CASE_PARTICIPANT_ID"].nunique())
print("Unique CASE_ID in analytic dataset:", df_analytic["CASE_ID"].nunique())
```

Original DataFrame shape: (248146, 52)

After subsetting to single-participant cases. (216252, 52)

After subsetting to primary charges with current sentences: (152900, 52)

Max unique charges per case-participant after subsetting: 1

Min unique charges per case-participant after subsetting: 1

Drop n/a judges: (152449, 52)

Number of rows in analytic dataset after subsetting: 135165

Sanity Check of Analytic Dataset:

Unique CASE_PARTICIPANT_ID in analytic dataset: 135165

Unique CASE_ID in analytic dataset: 135165