

Problem Set 3: Merging and Regular Expressions, Part 1

Background on the policy context: here, we're going to use two datasets to practice reshaping, merging, and regular expression patterns. Both datasets relate to the broader issue of which employers might be violating the rights of temporary guestworkers granted visas under the H-2A program. Here are some articles about potential exploitation of guestworkers by firms and inequality caused by minimal oversight:

- News media coverage of labor abuses of temporary guestworkers: <https://www.buzzfeednews.com/article/kenbensinger/the-pushovers>
- GAO report on labor abuses of temporary guestworkers: <https://www.gao.gov/products/gao-15-154>

The following datasets are located in `pset3_inputdata` (need to unzip):

- `jobs` : a dataset of guestworker jobs posted by many employers, some of whom have been debarred (banned) from the program for labor abuses; others not debarred
- `debar` : a dataset of employers who committed violations of labor regulations meant to protect temporary guestworkers

Here's a codebook explaining the variables in `jobs.csv` : https://web.archive.org/web/20231129014506/https://www.dol.gov/sites/dolgov/files/ETA/oflc/pdfs/H-2A_Record_Layout_FY2023_Q4.pdf

```
In [49]: ## helpful packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import random
import re
import os

## repeated printouts
from IPython.core.interactiveshell import InteractiveShell
InteractiveShell.ast_node_interactivity = "all"
```

Reshaping data (13 points total)

Load the following dataset stored in `pset3_inputdata : debar.csv`

This represents employers temporarily banned from hiring workers (debar.csv); call this `debar`

View the head()

```
In [50]: ## your code here
debar = pd.read_csv("pset3_inputdata/debar.csv")

debar.head()
```

	Name	City, State	Violation	Duration	Start date	End date
0	J&J Harvesting	Leads, ND	Failure to respond to audit (partial response)	2 years	1/19/2014	1/18/2016
1	Stahlman Apiaries, Inc	Selby, SD	Failure to respond to audit (partial response)	1 year	2/19/2015	2/14/2016
2	Trust Nursery	Pulaski, NY	Failure to respond to audit (partial response)	1 year	3/21/2014	3/20/2015
3	Anton Fertilizer Inc.	Dighton, KS	Failure to respond to audit (no response)	2 years	3/30/2014	3/29/2016
4	Great Plains Fluid Service, Inc.	Greensburg, KS	Failure to respond to audit (no response)	2 years	3/30/2014	3/29/2016

1.1 (1 point)

Print the number of rows in `debar` versus the number of unique employer names (`Name`). Is there one row per employer or multiple rows for some employers?

```
In [51]: ## your code here
num_rows = len(debar)
num_employers = debar["Name"].nunique()

print(f"The debar dataset has {num_rows} rows.")
print(f"There are {num_employers} unique employers in the dataset.")
print(f"Considering there is {num_rows - num_employers} more rows than unique employers, some employers must appear multiple times in the dataset.")
```

The debar dataset has 114 rows.

There are 98 unique employers in the dataset.

Considering there is 16 more rows than unique employers, some employers must appear multiple times in the dataset.

1.2 Investigating duplicated rows (2 points)

A. Create a new column in `debar` -- `is_repeated` -- that tells us whether an employer (`Name`) is repeated > 1 times

Hint: there are multiple ways to solve this but some possibilities to get the list of names that are repeated are:

- Using `value_counts()` on the `Name` variable and extracting the index from that value counts
- Using `groupby` to count the rows attached to one name

B. Print the rows where `is_repeated == True` and interpret

C. Subset to the rows where `is_repeated == True` and save that data as `mult_debar`. Print the `head()` and `shape`

```
In [52]: ## your code here
debar["is_repeated"] = debar["Name"].duplicated(keep=False)

mult_debar = debar[debar["is_repeated"] == True].copy()

mult_debar.head()
```

```
print(mult_debar.value_counts())
print(f"Shape: {mult_debar.shape}")
print(f"There are {mult_debar['Name'].nunique()} employers that appear multiple times in the dataset with a total of {len(mult_debar)} violations.")
```

Out[52]:

	Name	City, State		Violation	Duration	Start date	End date	is_repeated
6	Annabella Land & Cattle	Annabella, UT		Non Payment	1 year	5/9/2014	5/9/2015	True
7	Autumn Hill Orchard	Groton, MA	Failure to respond to audit (no response)		2 years	7/6/2014	7/5/2016	True
8	Caddo Creek Ranch, dba Paradise Ranch	Caddo, TX	Failure to respond to audit (partial response)		2 years	7/20/2014	7/19/2016	True
11	Loewen Harvesting LLC	Brownsville, TX	Failure to respond to audit (partial response)		1 year	8/20/2014	8/19/2015	True
12	Rollo Farm Labor Contractor	Miami, FL	Failure to respond to audit (no response)		2 years	8/23/2014	8/22/2016	True

Name	City, State	Violation	Duration	Start d
date	End date	is_repeated		
Annabella Land & Cattle	Annabella, UT	Non Payment	1 year	5/9/201
4 5/9/2015	True	1		
	Annabella, Utah	Non-payment	1 year	5/9/201
4 5/8/2015	True	1		
Autumn Hill Orchard	Groton, MA	Failure to respond to audit (no response)	2 years	7/6/201
4 7/5/2016	True	1		
		Impeding the Audit Process - Non- Response	2 years	7/6/201
4 7/5/2016	True	1		
Caddo Creek Ranch, dba Paradise Ranch	Caddo, TX	Failure to respond to audit (partial response)	2 years	7/20/20
14 7/19/2016	True	1		
	Caddo, Texas	Impeding the Audit Process - Partial- Response	2 years	7/20/20
14 7/19/2016	True	1		
Cisco Produce Inc.	Cairo, GA	Failure to respond to audit (no response)	2 years	12/10/2
014 12/9/2016	True	1		
		Impeding the Audit Process - Non- Response	2 years	12/10/2
015 12/9/2017	True	1		
Dove Creek Farms	Mount Vernon, TX	Failure to Respond to Audit Request	2 years	2/9/201
8 2/8/2020	True	1		
		Failure to respond to audit request	2 years	2/9/201
8 2/8/2018	True	1		
F&W Farms	Ingalls, KS	Failure to respond to audit (partial response)	2 years	12/10/2
014 12/9/2016	True	1		
		Impeding the Audit Process - Partial- Response	1 year	12/10/2
014 12/9/2015	True	1		
Loewen Harvesting LLC	Brownfield, TX	Impeding the Audit Process - Partial- Response	1 year	8/20/20
14 8/19/2015	True	1		
	Brownsville, TX	Failure to respond to audit (partial response)	1 year	8/20/20
14 8/19/2015	True	1		
Macky and Brad Farms	Plains, TX	Failure to respond to audit (no response)	1 year	2/13/20
15 2/12/2016	True	1		
		Impeding the Audit Process - Partial- Response	1 year	2/13/20
15 2/12/2016	True	1		
Maple Ridge Custom Services, LLC	Altheimer, AK	Failure to respond to audit (partial response)	2 years	11/16/2
014 11/15/2016	True	1		
	Altheimer, AR	Impeding the Audit Process - Partial- Response	1 year	11/16/2
014 11/15/2015	True	1		
Mark Duncan	Roosevelt, UT	Failure to respond to audit (no response)	2 years	11/16/2
014 11/15/2016	True	1		
		Impeding the Audit Process - Non- Response	2 years	11/16/2
014 11/15/2016	True	1		
Old Tree Farms/Verpaalen Custom Service	Volga, SD	WHD Debarment	3 years	12/11/2
014 12/10/2017	True	1		
		Wage Hour Debarment	3 years	12/1/20
14 12/1/2017	True	1		
Rollo Farm Labor Contractor	Miami, FL	Failure to respond to audit (no response)	2 years	8/23/20
14 8/22/2016	True	1		
		Impeding the Audit Process - Non- Response	2 years	8/23/20
14 8/22/2016	True	1		
SRT Farms	Morton, TX	Failure to respond to audit (no response)	2 years	11/16/2
014 11/15/2016	True	1		

```

014 11/15/2016 True      1           Impeding the Audit Process - Non- Response   2 years  11/16/2
Sharon Mathis                Tifton, GA          Failure to respond to audit (no response)  2 years  11/16/2
014 11/15/2016 True      1           Impeding the Audit Process - Non- Response   2 years  11/16/2
Turner Farms                  Healy, KS          Failure to comply with the employer's obligations to recruit U.S. workers 6 months  7/17/19
2/10/2020 True               1
2/10/2020 True               1           Lyons, Georgia      Failure to respond to audit request        2 years  9/27/20
Xavier Horne
17 9/26/2019 True            1           Non-payment of certification fee          1 year   6/16/20
16 6/15/2017 True            1
Name: count, dtype: int64
Shape: (32, 7)
There are 16 employers that appear multiple times in the dataset with a total of 32 violations.

```

1.3 Reshape mult_debar to wide to begin filtering out duplicates (4 points)

You want to separate out two cases:

- Cases where the repeat rows for one employer are due to duplicated data
- Cases where the repeat rows for one employer represent repeated violations for different issues

There are various ways to check duplicates in this data (eg converting `Violation` to lowercase; replacing spelled-out states with two-dig state codes)

We're going to use the simple rule of:

- A row is a duplicate if, within an employer (defined by Name + City, State), the Start date for each row's violation is the same

To begin to check this, reshape `mult_debar` to a wide dataframe (`mult_debar_wide`) with the following columns, treating the `Name` and `City, State` as the index for the pivot:

- Name
- City, State
- start_date_viol1
- start_date_viol2

Print the head and shape

In [53]: `## your code here`

```

# Create a violation number within each employer group
mult_debar[\"viol_num\"] = mult_debar.groupby([\"Name\", \"City, State\"]).cumcount() + 1

# Pivot to wide format with Start date as the values
mult_debar_wide = mult_debar.pivot(

```

```

    index=["Name", "City, State"],
    columns="viol_num",
    values="Start date"
).reset_index()

# Rename the columns to match required format
mult_debar_wide.columns = ["Name", "City, State", "start_date_viol1", "start_date_viol2"]

mult_debar_wide.head()
print(f"Shape: {mult_debar_wide.shape}")

```

Out[53]:

	Name	City, State	start_date_viol1	start_date_viol2
0	Annabella Land & Cattle	Annabella, UT	5/9/2014	NaN
1	Annabella Land & Cattle	Annabella, Utah	5/9/2014	NaN
2	Autumn Hill Orchard	Groton, MA	7/6/2014	7/6/2014
3	Caddo Creek Ranch, dba Paradise Ranch	Caddo, TX	7/20/2014	NaN
4	Caddo Creek Ranch, dba Paradise Ranch	Caddo, Texas	7/20/2014	NaN

Shape: (20, 4)

1.4 Filter out duplicates from original debar data (6 points)

A. Using `mult_debar_wide`, add a column `is_dup` that takes value of True for cases where `start_date_viol1 == start_date_viol2` marking the row as a duplicate

B. Going back to the original long-format data you loaded at the beginning -- `debar`

- Create a column in `debar` called `violnum` that numbers the violations for each employer (so the first violation for an employer is 1, the second is 2, etc)
- For employers where `is_dup == True` as indicated by your wide-format dataframe, only keep `violnum == 1`
- For all other employers (so `is_dup == False` and ones we didn't need to check duplicates for), keep all `violnum`
- Remove the `is_repeated` column from the `debar` data

Hint: you can complete part B without a for loop; `pd.concat` with axis = 0 (row binding) is one way

Call the resulting dataframe `debar_clean` and print the shape and # of unique employer names

Save the `debar_clean` as a csv file using `df.to_csv([filename], index=False)`. You will be using it in Problem Set 4.

```

In [ ]: # Part A: Add is_dup column to mult_debar_wide
mult_debar_wide["is_dup"] = (mult_debar_wide["start_date_viol1"] == mult_debar_wide["start_date_viol2"])
print(mult_debar_wide["is_dup"].value_counts())

# Part B: Filter duplicates from original debar data

# Create violnum column

```

```

print(f"original debar shape: {debar.shape}")

debar = debar.sort_values(["Name", "City, State"]).copy()
debar["violnum"] = debar.groupby(["Name", "City, State"]).cumcount() + 1

# Get list of duplicate employers (Name + City, State) from the wide frame
dup_employers = mult_debar_wide.loc[mult_debar_wide["is_dup"], ["Name", "City, State"]].drop_duplicates()

# Mark rows in debar that belong to duplicate employers
debar_flagged = debar.merge(dup_employers, on=["Name", "City, State"], how="left", indicator=True)

# Keep only violnum==1 for duplicate employers; keep all rows for others
debar_dup = debar_flagged.loc[(debar_flagged["_merge"] == "both") & (debar_flagged["violnum"] == 1)]
debar_not_dup = debar_flagged.loc[debar_flagged["_merge"] == "left_only"]

# Combine and drop helper columns (is_repeated, violnum, _merge)
debar_clean = pd.concat([debar_dup, debar_not_dup], axis=0).drop(columns=["is_repeated", "violnum", "_merge"])

# Print shape and # of unique employer names
print(f"debar_clean shape: {debar_clean.shape}")
print(f"# unique employer names: {debar_clean['Name'].nunique()}")

# Save to csv
debar_clean.to_csv("debar_clean.csv", index=False)

```

```

is_dup
False    11
True     9
Name: count, dtype: int64
original debar shape: (114, 10)
debar_clean shape: (105, 8)
# unique employer names: 98

```

```

In [58]: # Inspect which employers are being deduplicated
print(f"Employers with duplicate violations (removed to keep only one): {len(dup_employers)}")
print(dup_employers)

# Show all rows for these employers in the original debar
dup_names = dup_employers[["Name", "City, State"]].values.tolist()
print("Rows in debar for deduplicated employers:")
for name, city_state in dup_names:
    display(debar[(debar["Name"] == name) & (debar["City, State"] == city_state)])

```

Employers with duplicate violations (removed to keep only one): 9

	Name	City, State
2	Autumn Hill Orchard	Groton, MA
6	Dove Creek Farms	Mount Vernon, TX
7	F&W Farms	Ingalls, KS
10	Macky and Brad Farms	Plains, TX
13	Mark Duncan	Roosevelt, UT
15	Rollo Farm Labor Contractor	Miami, FL
16	SRT Farms	Morton, TX
17	Sharon Mathis	Tifton, GA
18	Turner Farms	Healy, KS

Rows in debar for deduplicated employers:

	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
7	Autumn Hill Orchard	Groton, MA	Failure to respond to audit (no response)	2 years	2014-07-06	7/5/2016	True	1.0	2014	MA
29	Autumn Hill Orchard	Groton, MA	Impeding the Audit Process – Non- Response	2 years	2014-07-06	7/5/2016	True	2.0	2014	MA

	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
103	Dove Creek Farms	Mount Vernon, TX	Failure to respond to audit request	2 years	2018-02-09	2/8/2018	True	1.0	2018	TX
109	Dove Creek Farms	Mount Vernon, TX	Failure to Respond to Audit Request	2 years	2018-02-09	2/8/2020	True	2.0	2018	TX

	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
18	F&W Farms	Ingalls, KS	Failure to respond to audit (partial response)	2 years	2014-12-10	12/9/2016	True	1.0	2014	KS
58	F&W Farms	Ingalls, KS	Impeding the Audit Process – Partial- Response	1 year	2014-12-10	12/9/2015	True	2.0	2014	KS

	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
31	Macky and Brad Farms	Plains, TX	Failure to respond to audit (no response)	1 year	2015-02-13	2/12/2016	True	1.0	2015	TX
55	Macky and Brad Farms	Plains, TX	Impeding the Audit Process – Partial- Response	1 year	2015-02-13	2/12/2016	True	2.0	2015	TX

	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
16	Mark Duncan	Roosevelt, UT	Failure to respond to audit (no response)	2 years	2014-11-16	11/15/2016	True	1.0	2014	UT
60	Mark Duncan	Roosevelt, UT	Impeding the Audit Process – Non- Response	2 years	2014-11-16	11/15/2016	True	2.0	2014	UT

	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
12	Rollo Farm Labor Contractor	Miami, FL	Failure to respond to audit (no response)	2 years	2014-08-23	8/22/2016	True	1.0	2014	FL
24	Rollo Farm Labor Contractor	Miami, FL	Impeding the Audit Process – Non- Response	2 years	2014-08-23	8/22/2016	True	2.0	2014	FL

	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
15	SRT Farms	Morton, TX	Failure to respond to audit (no response)	2 years	2014-11-16	11/15/2016	True	1.0	2014	TX
61	SRT Farms	Morton, TX	Impeding the Audit Process – Non- Response	2 years	2014-11-16	11/15/2016	True	2.0	2014	TX
	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
14	Sharon Mathis	Tifton, GA	Failure to respond to audit (no response)	2 years	2014-11-16	11/15/2016	True	1.0	2014	GA
62	Sharon Mathis	Tifton, GA	Impeding the Audit Process – Non- Response	2 years	2014-11-16	11/15/2016	True	2.0	2014	GA
	Name	City, State	Violation	Duration	Start date	End date	is_repeated	violnum	Year	State
106	Turner Farms	Healy, KS	Failure to comply with the employer's obligati...	6 months	2019-07-17	2/10/2020	True	1.0	2019	KS
111	Turner Farms	Healy, KS	Failure to comply with the employer's obligati...	7 months	2019-07-17	2/10/20	True	2.0	2019	KS

2. Optional extra credit 2 (up to 3 points)

- For 1 point extra credit, create a visualization with 1+ of the existing fields in either the raw `jobs` or `debar` data. We'll be showing cool visualizations in class so use your imagination! Options could include visualizing between-state or over-time variation
- For 3 points extra credit instead, geocode the employer addresses in `jobs` and plot the addresses of jobs as points overlaid on top of a map of Georgia
 - Note:** this extra credit involves Googling since we have not yet covered spatial data.
 - For discussion of how to geocode addresses -> lat/long, see: <https://www.natasshaselvaraj.com/a-step-by-step-guide-on-geocoding-in-python/>
 - For discussion of plotting lat/long dots against a map, see this discussion of geopandas: <https://towardsdatascience.com/plotting-maps-with-geopandas-428c97295a73>
 - Relevant columns include `EMPLOYER_ADDRESS_1`
 - The geocoding might have a long runtime so feel free to implement it in a separate .py script that you submit alongside your notebook and to just read in the geocoded data

```
In [ ]: # Convert Start date to datetime and extract year
debar["Start date"] = pd.to_datetime(debar["Start date"], format="mixed")
debar["Year"] = debar["Start date"].dt.year

# Extract state from "City, State" column
debar["State"] = debar["City, State"].str.extract(r',\s*([A-Z]{2})$')

# Create a figure with two subplots
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

# Violations over time by year
violations_by_year = debar.groupby("Year").size()
_ = axes[0].bar(violations_by_year.index, violations_by_year.values, color="steelblue", edgecolor="black")
_ = axes[0].set_xlabel("Year")
```


```

_ = axes[0].set_ylabel("Number of Violations")
_ = axes[0].set_title("H-2A Program Violations Over Time")
_ = axes[0].set_xticks(violations_by_year.index)

# Top 10 states by number of violations
_ = violations_by_state = debar["State"].value_counts().head(10)
_ = axes[1].barh(violations_by_state.index[::-1], violations_by_state.values[::-1], color="coral", edgecolor="black")
_ = axes[1].set_xlabel("Number of Violations")
_ = axes[1].set_ylabel("State")
_ = axes[1].set_title("Top 10 States with H-2A Violations")
plt.tight_layout()
plt.show()

print(f"\nSummary: The data shows {len(debar)} total violations across {debar['State'].nunique()} states from {debar['Year'].min()} to {debar['Year'].max()}.")

```



Summary: The data shows 114 total violations across 23 states from 2013 to 2019.

In [38]: %run georgia_jobs_map.py

Number of Georgia employers: 74
Unique addresses to geocode: 68
No cache found, starting fresh...

Full address failed, trying simple address: Lyons, GA 30436
Full address failed, trying simple address: Tifton, GA 31793
Full address failed, trying simple address: Wray, GA 31798
Full address failed, trying simple address: Lyons, GA 30436
Full address failed, trying simple address: Cobbtown, GA 30420
Full address failed, trying simple address: Adel, GA 31620
Full address failed, trying simple address: Parrott, GA 39877

Geocoded 10/68 addresses...

Full address failed, trying simple address: Leesburg, GA 31763
Full address failed, trying simple address: Omega, GA 31775
Full address failed, trying simple address: Lyons, GA 30436
Full address failed, trying simple address: Lyons, GA 30436
Full address failed, trying simple address: Cobbtown, GA 30420
Full address failed, trying simple address: Norman Park, GA 31771
Full address failed, trying simple address: Pitts, GA 31072

Geocoded 20/68 addresses...

Full address failed, trying simple address: Cordele, GA 31015
Full address failed, trying simple address: Lake Park, GA 31636
Full address failed, trying simple address: Alma, GA 31510
Full address failed, trying simple address: Chatsworth, GA 30705
Full address failed, trying simple address: Reidsville, GA 30453
Full address failed, trying simple address: Baxley, GA 31513
Full address failed, trying simple address: Alma, GA 31510
Full address failed, trying simple address: Baxley, GA 31513
Full address failed, trying simple address: Moultrie, GA 31768

Geocoded 30/68 addresses...

Full address failed, trying simple address: Tifton, GA 31793
Full address failed, trying simple address: Lyons, GA 30436
Full address failed, trying simple address: Tifton, GA 31794
Full address failed, trying simple address: Alma, GA 31510
Full address failed, trying simple address: Hoschton, GA 30548
Full address failed, trying simple address: Blairsville, GA 30514
Full address failed, trying simple address: Washington, GA 30673
Full address failed, trying simple address: Leslie, GA 31764

Geocoded 40/68 addresses...

Full address failed, trying simple address: Baxley, GA 31513
Full address failed, trying simple address: Norman Park, GA 31771
Full address failed, trying simple address: Reidsville, GA 30453
Full address failed, trying simple address: Claxton, GA 30417
Full address failed, trying simple address: Sycamore, GA 31790
Full address failed, trying simple address: Dearing, GA 30808
Full address failed, trying simple address: Fort Valley, GA 31030
Full address failed, trying simple address: Milner, GA 30257


Geocoded 50/68 addresses...


Full address failed, trying simple address: Broxton, GA 31519
Full address failed, trying simple address: Alto, GA 30510
Full address failed, trying simple address: Crawford, GA 30630
Full address failed, trying simple address: Patterson, GA 31557
Full address failed, trying simple address: Rochelle, GA 31079

```
Full address failed, trying simple address: Lyons, GA 30436
Geocoded 60/68 addresses...
Full address failed, trying simple address: Wray, GA 31798
Full address failed, trying simple address: Alma, GA 31510
Full address failed, trying simple address: Leesburg, GA 31763
Full address failed, trying simple address: Whitesburg, GA 30185
Full address failed, trying simple address: Adrian, GA 31002
```

```
Successfully geocoded: 68 / 68 addresses
Results saved to geocoded_addresses_cache.csv
Jobs with valid coordinates: 74
```

H-2A Guestworker Job Locations in Georgia





Plotted 74 employer job locations on Georgia map.

In []: